


HBase

Topics at a glance

- Hadoop and its data access limitations.
- Why Hbase?
- Hbase and its importance in Hadoop frame work.
- History and Architecture of Hbase.
- Hbase components and their responsibilities.
- Hbase data storage model.
- Advantages and disadvantages of Hbase.
- Conclusion of the session



Is there a way **Hadoop** can
find Customer Addresses
from this pile of
Unstructured Data



Why Hbase???

With the evolution of the internet web application scope was increased -

- Huge volumes of structured and semi-structured data started getting generated.
- Semi-structured data (emails, JSON, XML, and .csv files and exe files)
- Loads of semi-structured data was created across the globe.
- So storing and processing of this data became a major challenge.

Solution



Hadoop and its limitations

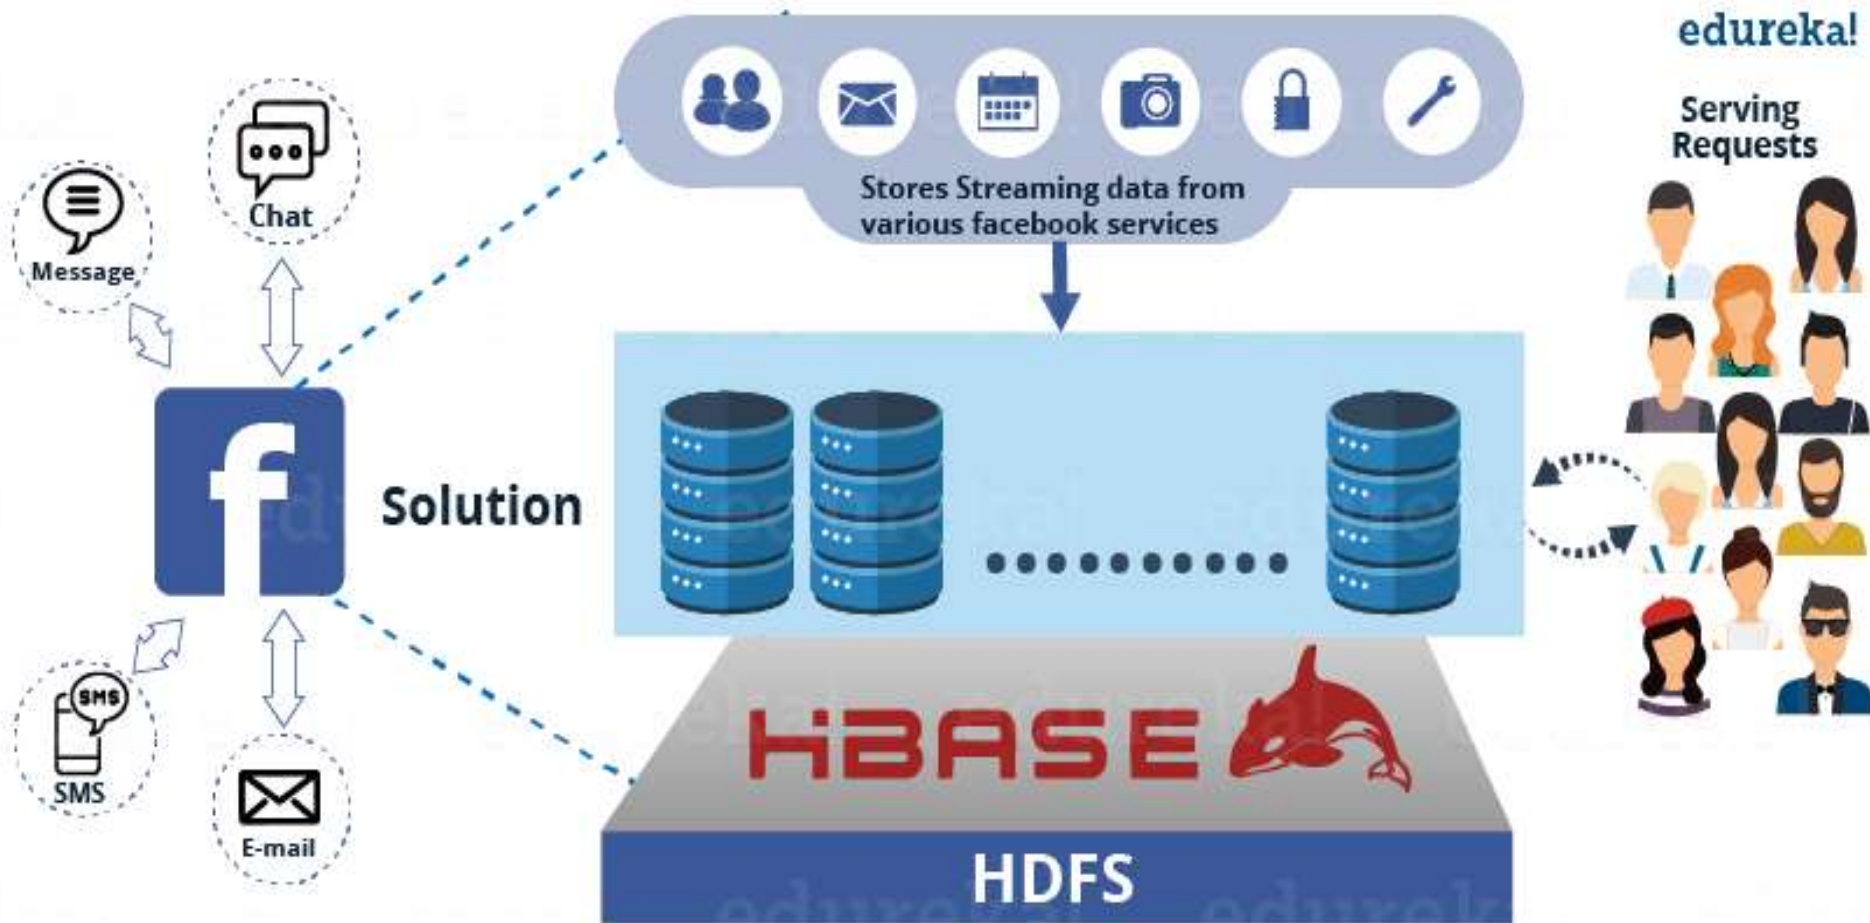
- Hadoop can perform only batch processing, and data will be accessed only in a sequential manner.
- So if needed to access any data randomly then need a new access methodology.

✓ New program tool added in Hadoop framework to provide the random access to user.

- HBase
- Cassandra,
- CouchDB,
- Dynamo and MongoDB

HBase

- HBase is an open-source NoSQL database and Part of the Hadoop framework.
- Similar to Google's big table. Initially, it was Google Big Table, afterward; it was renamed as HBase .
- Hbase is primarily written in **Java** and needed for **real-time Big Data** applications.
- HBase is a distributed column-oriented **non-relational** database management system that runs on top of Hadoop Distributed File System (HDFS).
- HBase is a **column-oriented database** and the tables in it are sorted by row.
- The table schema defines only column families, which are the key value pairs.
- It uses log storage with the Write-Ahead Logs (WAL).
- It supports fast random access and heavy writing competency.



How is HBase different from other NoSQL models

- HBase stores data in the form of key/value pairs in a columnar model. In this model, all the columns are grouped together as Column families.
- HBase on top of Hadoop will increase the throughput and performance of distributed cluster set up.
- Provides faster random reads and writes operations.

HBase	RDBMS
<ul style="list-style-type: none">• HBase works well with structured and semi-structured data• Hbase does not have a fixed schema. Here, only column families are defined• It can have <u>denormalized data</u>• It is built for wide tables that can be scaled horizontally	<ul style="list-style-type: none">• RDBMS has a fixed schema which describes the structure of the tables• Works well with structured data• RDBMS can store only normalized data• It is built for thin tables that is hard to scale

Features of Hbase

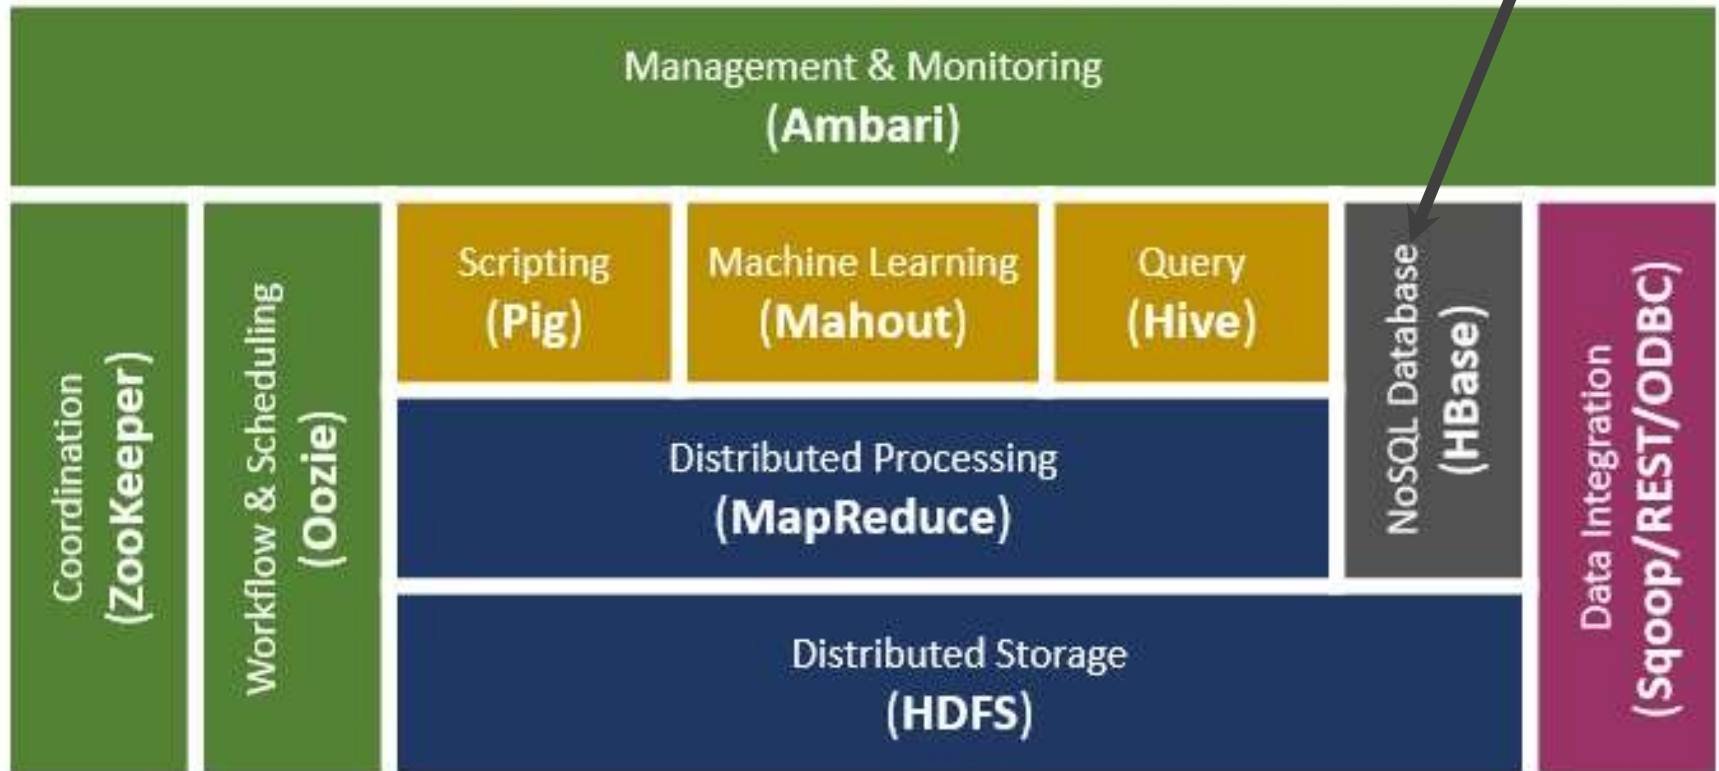
- **Horizontally scalable:** Can add any number of columns anytime.
- **Automatic Failover:** Allows a system administrator to automatically switch data handling to a standby system in the event of system compromise/failure.
- **Integrations with Map/Reduce framework:** All the commands and java codes internally implement Map/ Reduce to do the task and it is built over Hadoop Distributed File System.
- It doesn't enforce relationships within your data.
- It is designed to run on a cluster of computers, built using commodity hardware.
- HBase is built for low latency operations

History of HBase


- In Nov 2006, Google released the paper on BigTable.
- Feb 2007, Initial HBase prototype was created as a Hadoop contribution.
- Oct 2007, The first usable HBase along with Hadoop 0.15.0 was released.
- Jan 2008, HBase became the sub project of Hadoop.
- Oct 2008, HBase 0.18.1 was released.
- Jan 2009, HBase 0.19.0 was released.
- Sept 2009, HBase 0.20.0 was released.
- May 2010, HBase became Apache top-level project.

HBase existence in the Hadoop Ecosystem

Apache Hadoop Ecosystem



HBase Table – To store data



HBase: Keys and Column Families

Each record is divided into Column Families

Each row has a Key

PERSON TABLE					
row key	personal_data		demographic		...
PersonID	Name	Address	BirthDate	Gender	...
1	H. Houdini	Budapest, Hungary	1926-10-31	M	
2	D. Copper	New Jersey, USA	1956-09-16	M	
3	Merlin	Stonehenge, England	1136-12-03	F	
...	
500,000,000	F. Cadillac	Nevada, USA	1964-01-07	M	

Figure 2 - Census Data in Column Families

Each column family consists of one or more Columns

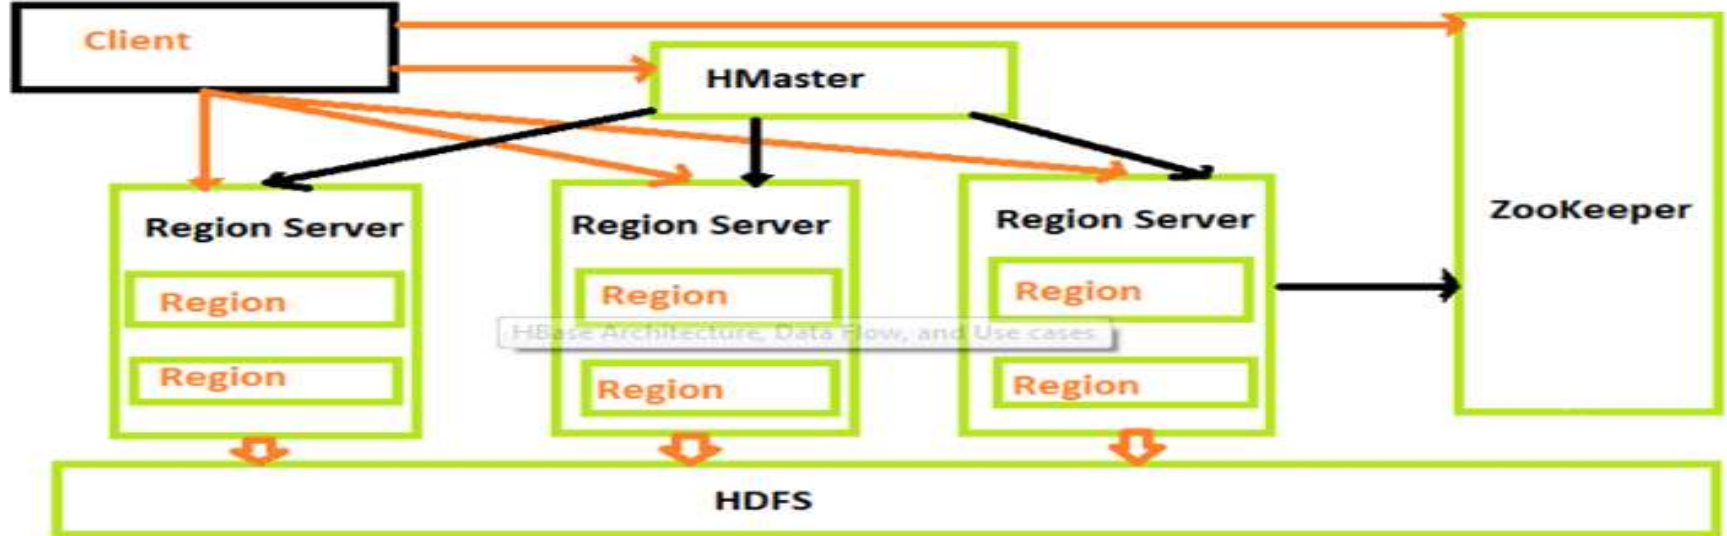
Example- Storage Mechanism in HBase

- HBase is a column-oriented database.
- Data is store in form of table.

Row Key		Column Family			Column Qualifiers
Row Key		Customers		Products	
Customer ID	Customer Name	City & Country	Product Name	Price	Cell
1	Sam Smith	California, US	Mike	\$500	
2	Arijit Singh	Goa, India	Speakers	\$1000	
3	Ellie Goulding	London, UK	Headphones	\$800	
4	Wiz Khalifa	North Dakota, US	Guitar	\$2500	

Figure: HBase Table

HBase Architecture

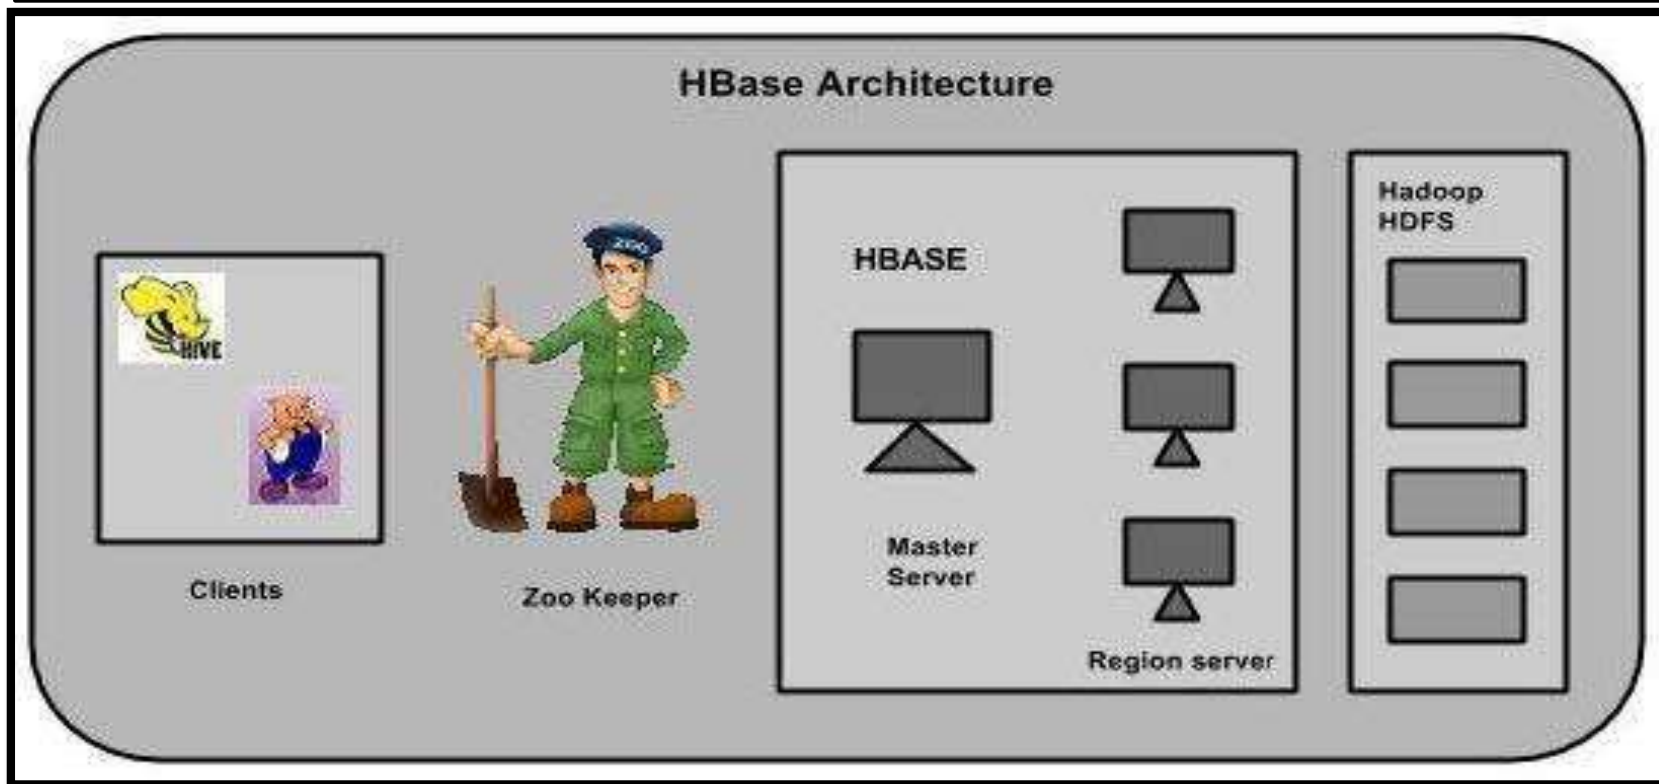


- ❖ Apache Zookeeper monitors the system.
- ❖ HBase Master assigns regions and load balancing.
- ❖ The Region server serves data to read and write.
- ❖ The Region Server is all the different computers in the Hadoop cluster.
- ❖ It consists of Region, HLog, Store, Memory Store, and different files.
- ❖ All this is a part of the HDFS storage system.

HBase - Components

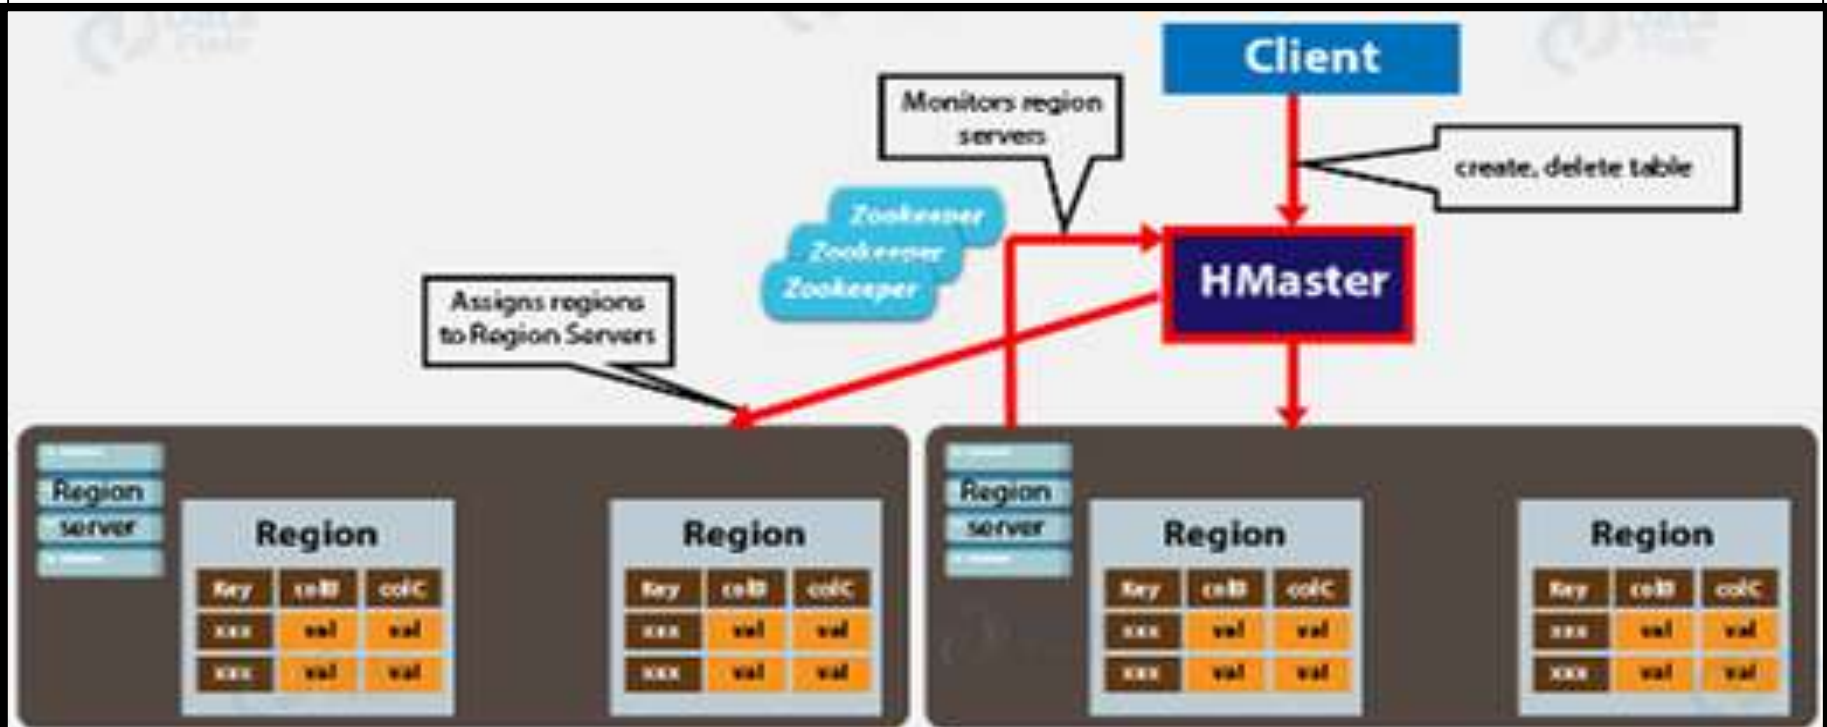
HBase has three major components -

1. Master servers.
2. Region server
3. Zookeeper



1. HMaster

- **HMaster** in HBase is the implementation of a Master server in HBase.
- It acts as a monitoring agent to monitor all Region Server instances present in the cluster and acts as an interface for all the metadata changes.
- In a distributed cluster environment, Master runs on **Name Node**.



a. Coordinating the region servers as following -

Assigns Regions on startup.

Recovery and load balancing.

Monitors all RegionServer instances in the HBase Cluster.

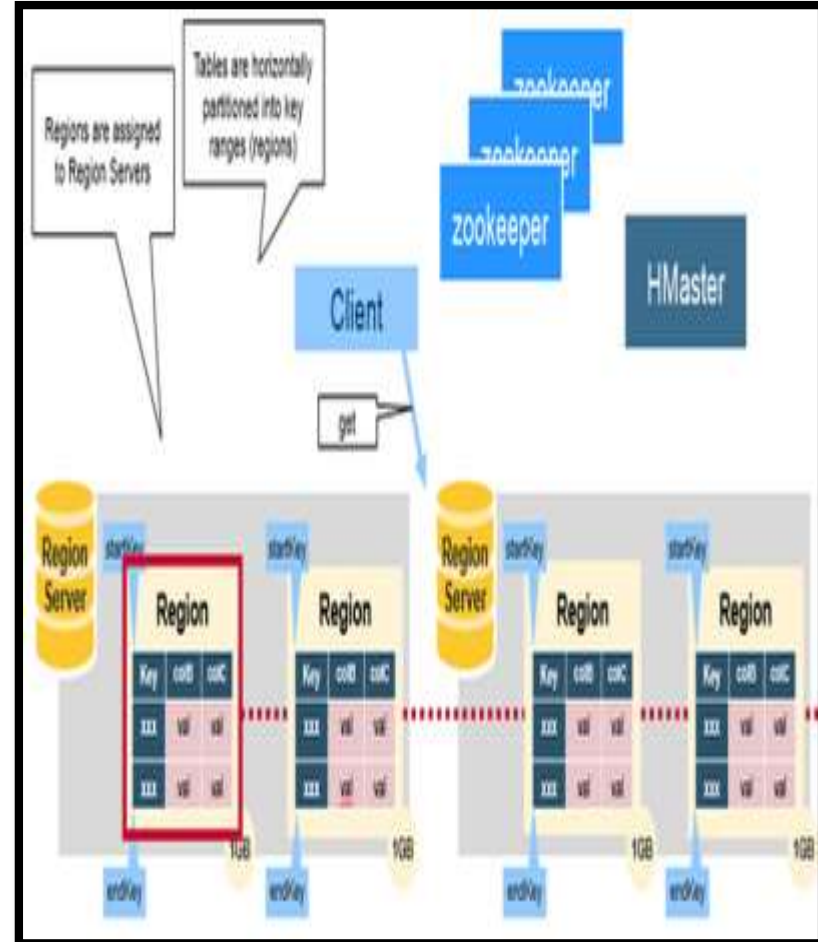
b. Admin functions

When a client wants to change any schema or change in Metadata operations, HMaster takes responsibility for these operations as follow-

- Table (create able, removeTable, enable, disable)
- ColumnFamily (add Column, modify Column)
- Region (move, assign)

2. Regions & Regions Server

- **Table** is **split** According to **rowkey** Scope **horizontal** to several **region**. (**Start to end key**)
- After split rows are called region and these Regions are assigned to certain nodes in the cluster for management is called **Region Server**.
- They are responsible for processing data read and write requests.
- Each Region Server can manage about 1000 regions.
- HRegion Server is the Region Server implementation.
- Responsible for serving and managing regions or data that is present in a distributed cluster.
- Region servers run on Data Nodes present in the Hadoop cluster.



How Regions splits

Sr.No	Personal Details		Education_ Details		Job details	
Emp_id	Name	Age	Graduate	Percent age	Company Name	Designation
1	Amit	25	Bsc	76	HCL	Project Lead
2	Sumit	30	BTech	80	TCS	Project Manager
3	Varsha	35	MTech	75	Wipro	Project Engineer

HDFS

Region
Server

1	Amit	25	Bsc	76	HCL	Project Lead
2	Sumit	30	BTech	80	TCS	Project Manager
3	Varsha	35	MTech	75	Wipro	Project Engineer

- When HBase Region Server receives writes and read requests from the client, it assigns the request to a specific region, where the actual column family resides.
- Client can directly contact with HRegion servers, there is no need of HMaster mandatory permission to the client regarding communication with HRegion servers.
- The client requires HMaster help when operations related to metadata and schema changes are required.

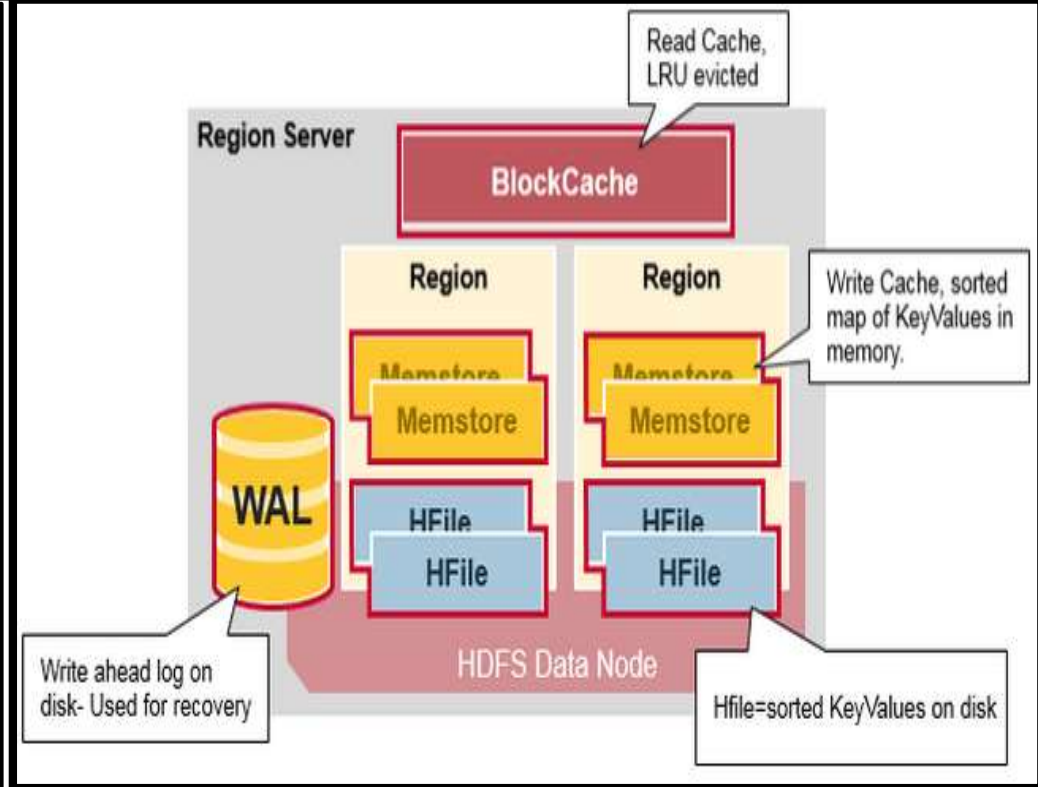
❖ HMaster can get into contact with multiple HRegion servers and performs the following functions.

- Hosting and managing regions
- Splitting regions automatically
- Handling read and writes requests
- Communicating with the client directly

- Region Server runs on HDFS DataNode and Responsible for processing data read and write requests.
- If Client required any data then he will directly interacts with Region Server.
- Regions are tables that are split up and spread across the region servers.

Components of Region server :

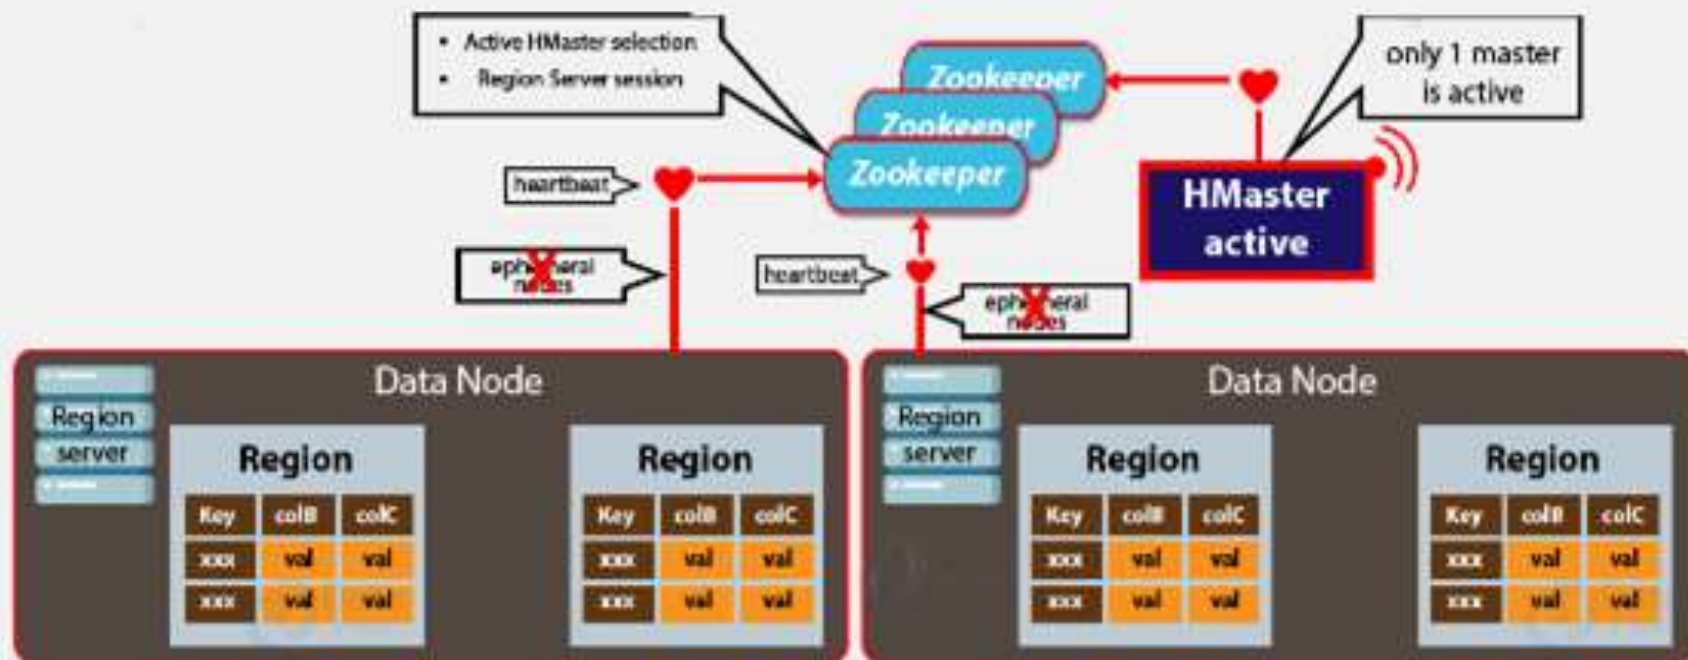
- **WAL (Write Ahead Log)** is a file on a distributed file system for **Storing new data**
- **Block Cache** - This is the read cache. **Memory** The most frequently accessed data is stored in the LRU (Least Current Used) cache.
- **MemStore** - This is the write cache, in **Memory**.
- **Hfile** -Store HBase data on hard disk (HDFS).



3.Hbase- Zookeeper

- Hbase use Zookeeper to coordinate shared state information for members of distributed systems.
- Active HMaster and Region servers, connects with a session to Zookeeper.
- For active sessions ZooKeeper maintains ephemeral nodes by using heartbeats.
- Zookeeper maintains which servers are healthily available and notifies them when the server fails.
- **Ephemeral nodes** mean znodes which exist as long as the session which created the znode is active and then znode is deleted when the session ends.
- Zookeeper uses a consistency protocol to ensure the consistency of the distributed state.
- Each Region Server in HBase Architecture produces an ephemeral node. Further, to discover available region servers and HMaster shall monitors these nodes.
- Active HMaster sends heartbeats to Zookeeper.

How the Components Work Together



1. Active HMaster

2. Inactive Hmaster

❖ HBase META Table

- META Table is a special HBase Catalog Table. Basically, it holds the location of the regions in the HBase Cluster.
- It keeps a list of all Regions in the system.
- Structure of the .META. table is as follows:
- **Key:** region start key, region id
- **Values:** RegionServer

- **Tables:** Data is stored in a table format in Hbase.
- **Row Key:** Row keys are used to search records which make searches fast.
- **Column Families:** Various columns are combined in a column family. These column families are stored together which makes the searching process faster because data belonging to same column family can be accessed together in a single seek.
- **Column Qualifiers:** Each column's name is known as its column qualifier.
- **Cell:** Data is stored in cells.
- **Timestamp:** Timestamp is a combination of date and time. Whenever data is stored, it is stored with its timestamp.

HDFS vs. HBase

HDFS

HBase

HDFS is a Java-based file system utilized for storing large data sets. HBase is a Java based No-SQL database.

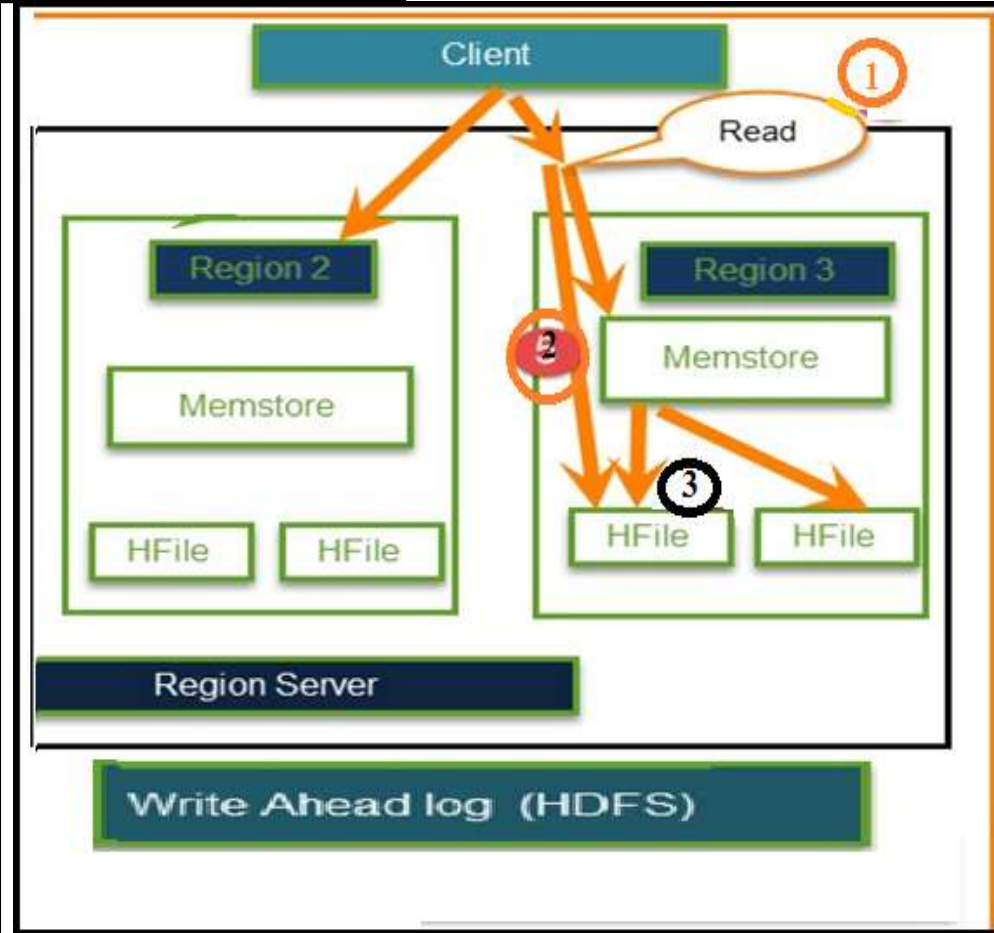
HDFS has a rigid architecture that does not allow changes. It doesn't facilitate dynamic storage. HBase allows for dynamic changes and can be utilized for standalone applications.

HDFS is ideally suited for write-once and read-many times use cases. HBase is ideally suited for random write and read of data that is stored in HDFS.

HBase - Read

- A Read against HBase must be reconciled between the HFiles, MemStore & BLOCKCACHE.
- The Block Cache is designed to keep frequently accessed data from the HFiles in memory so as to avoid disk reads.
- Each column family has its own Block Cache.

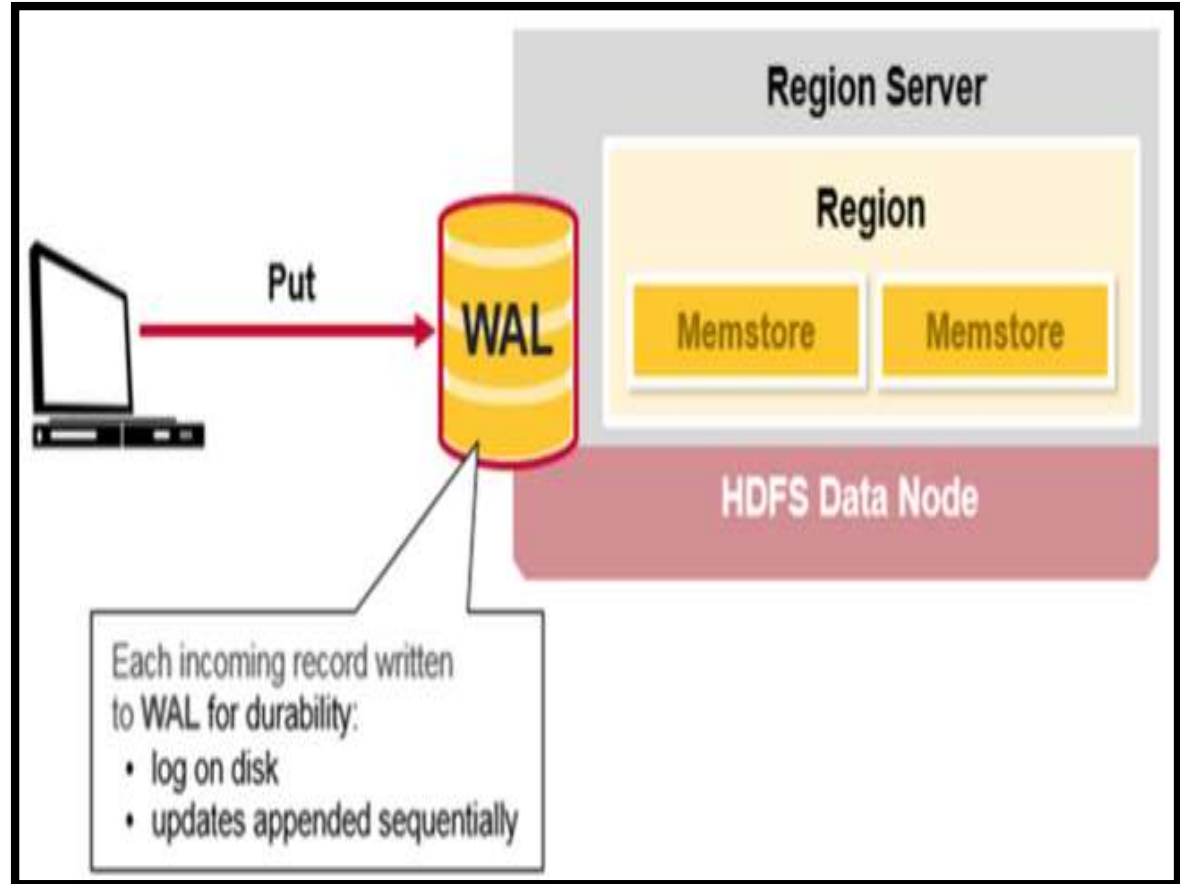
Block: It is the smallest indexed unit of data and is the smallest unit of data that can be read from disk. default size 64KB.



Hbase - Write

When a write is made, by default, it goes into two places:

- write-ahead log (WAL), Hlog.
- in-memory write buffer, MemStore.



- ❖ Hbase designed to store Denormalized Data.
- ❖ Hbase Supports Automatic Partitioning
- ❖ **Strong consistency model**– All readers will see same value, while a write returns.
- ❖ **Scales automatically**
 - While data grows too large, Regions splits automatically.
 - To spread and replicate data, it uses HDFS.
- ❖ **Built-in recovery** – It uses Write Ahead Log for recovery.
- ❖ **Integrated with Hadoop**
- ❖ Hbase is **schema-less, no data model has been defined.**
- ❖ Hbase has the ability to perform Random read and write operations.
- ❖ Hbase provides **data replication** across clusters for higher availability.
- ❖ Feature **random access** (internal hash table) to stores data in HDFS files for faster lookups/searching.

Disadvantages of HBase

- **Single point of failure** - If HMaster goes down, complete cluster will be fail and no work/task will be performed.
- Cannot perform functions like SQL and doesn't support SQL structure.
- Does not contain any query optimizer
- Does not support for **transaction**.
- **Business continuity reliability**
 - Write Ahead Log replay very slow.
 - Also, a slow complex crash recovery.
- **Joining** and **normalization** is very **difficult** to perform.
- Very difficult to store large **binary data**.

Real Time Example of HBase-Facebook

How Facebook use Hbase to store user data



User Account ID	Account Type (Personal/ Business)	Type of Contents Posted	Posted for (Public/ Private)	Time Stamp	Violating community standards	Last Login Activity Time of Account
Rahul3@fb	Personal	Image + Text	Public	20/08/2022 08.45.00.PM	No	05.30 PM
ABZ@fb	Business	Text +Video	Public	25/08/2022 06.45.00.PM	No	08.30 PM
Tarun@fb	Personal	Text +Video	Public	25/08/2022 06.45.00.PM	Yes	08.30 PM

User Account ID	Account Type (Personal/ Business)	Type of Contents Posted	Action Required	Account Suspended	Account Blocked	Remarks
Tarun@fb	Personal	Text +Video Is a hated contents	Yes	For a period 2 week etc.	Yes	Violating Community standard



Any Query??

